

Orange Juice Sales at Wasatch Grocery Chain

Identification of Significant Predictor Variables and Predictive Modelling of Customer
Preference in Minute Maid Sales

Chris Gearheart and Chris Porter

2022-11-15

Introduction

Wasatch Grocery Chain (WGC) is a regional grocery chain operating in the Intermountain West of the US. WGC sells two brands of orange juice in its stores, Citrus Hill (CH) and Minute Maid (MM) of which MM is the more profitable to the company. This report will identify what customer factors within available data contribute to purchase of MM over CH, as well as to what degree these factors influence customer choice. In addition, a predictive model is created that will allow the Sales Department to identify other customers within our customer base that are more likely to purchase Minute Maid brand orange juice, thus driving profitability across the company.

```
set.seed(1234)
df <- read.csv(url("http://data.mishra.us/files/project/OJ_data.csv"))
df[-1] <- lapply(df[-1], as.numeric)
df$Purchase <- as.factor(df$Purchase)
purchase_testtrain <- initial_split(df, prop = 0.75, strata = Purchase)
train <- training(purchase_testtrain)
test <- testing(purchase_testtrain)
```

Available Data

The data set used in this report contains 13 possible predictor variables as well as 1 outcome variable, Purchase, which records whether or not a customer purchased MM. There are a total of 1070 observations in the data set. The data set was further partitioned into a **training** data set, containing 801 observations, and a validation **testing** data set containing 269 observations.

Methods

Results

```
set.seed(1234)
recipe_oj <- recipe(Purchase ~ ., train)

model_oj_bt <- boost_tree(trees = tune(), tree_depth = tune(), learn_rate = tune()) %>%
  set_engine('xgboost', verbosity = 0) %>%
  set_mode('classification')

hyperparameter_grid <- grid_regular(trees(), tree_depth(), learn_rate(), levels = 5)

purchase_folds <- vfold_cv(train, v=4) # 4-fold Cross validation

oj_workflow <- workflow() %>% add_model(model_oj_bt) %>% add_recipe(recipe_oj) #Set Workflow

# Tune Hyper-parameters
oj_tune <- oj_workflow %>% tune_grid(resamples = purchase_folds,
                                     grid = hyperparameter_grid,
                                     metrics = metric_set(accuracy))

best_bt_model <- oj_tune %>% select_best('accuracy') #Select best Hyper-parameters from grid

best_bt_model


## # A tibble: 1 x 4
##   trees tree_depth learn_rate .config
##   <int>      <int>      <dbl> <chr>
## 1  1000          1        0.1 Preprocessor1_Model023
```

```

oj_final_workflow <- oj_workflow %>% finalize_workflow(best_bt_model) # Create Final Workflow based upon
final_fit <- oj_final_workflow %>% last_fit(split = purchase_testtrain) # Final Fit Model
final_fit %>% collect_metrics()

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>      <dbl> <chr>
## 1 accuracy binary      0.799 Preprocessor1_Model1
## 2 roc_auc  binary      0.892 Preprocessor1_Model1
oj_final_workflow %>% fit(data = train) %>% extract_fit_parsnip() %>% vip(geom = 'col') #Plot most important

```



```

vi_values <- oj_final_workflow %>% fit(data = train) %>% extract_fit_parsnip() %>% vi()
vi_gt_1 <- vi_values %>% filter(Importance >= 0.01)
vi_gt_1

```

```

## # A tibble: 5 x 2
##   Variable      Importance
##   <chr>         <dbl>
## 1 LoyalCH      0.824
## 2 PriceDiff    0.104
## 3 DiscCH       0.0239
## 4 ListPriceDiff 0.0232
## 5 SalePriceMM  0.0104

```

the most important variable is LoyalCH with a 82.3503048%

```
model_fitted <- oj_final_workflow %>% fit(data = train)


explainer_rf <- explain_tidymodels(model_fitted,
                                   data = train[,-1],
                                   y = train$Purchase,
                                   type = "pdp", verbose = FALSE)

pdp_LoyalCH <- model_profile(explainer_rf,
                             variables = "LoyalCH",
                             N=NULL)
pdp_PriceDiff <- model_profile(explainer_rf,
                              variables = "PriceDiff",
                              N=NULL)
pdp_DiscCH <- model_profile(explainer_rf,
                            variables = "DiscCH",
                            N=NULL)
pdp_ListPriceDiff <- model_profile(explainer_rf,
                                   variables = "ListPriceDiff",
                                   N=NULL)
pdp_SalePriceMM <- model_profile(explainer_rf,
                                 variables = "SalePriceMM",
                                 N=NULL)
pdp_DiscMM <- model_profile(explainer_rf,
                           variables = "DiscMM",
                           N=NULL)

plot(pdp_LoyalCH)
```

Partial Dependence profile


Created for the workflow model



```
plot(pdp_PriceDiff)
```

Partial Dependence profile

Created for the workflow model
PriceDiff



```
#plot(pdp_DiscCH)
#plot(pdp_ListPriceDiff)
#plot(pdp_SalePriceMM)
#plot(pdp_DiscMM)
```

Conclusions and Recommendations

Brand

Sales

Appendix 1: Data Characteristics

```
summary(df)
```

```
## Purchase      PriceCH      PriceMM      DiscCH      DiscMM
## 0:417   Min.   :1.690   Min.   :1.690   Min.   :0.00000   Min.   :0.0000
## 1:653   1st Qu.:1.790   1st Qu.:1.990   1st Qu.:0.00000   1st Qu.:0.0000
##        Median :1.860   Median :2.090   Median :0.00000   Median :0.0000
##        Mean   :1.867   Mean   :2.085   Mean   :0.05186   Mean   :0.1234
##        3rd Qu.:1.990   3rd Qu.:2.180   3rd Qu.:0.00000   3rd Qu.:0.2300
##        Max.   :2.090   Max.   :2.290   Max.   :0.50000   Max.   :0.8000
## SpecialCH    SpecialMM    LoyalCH    SalePriceMM
## Min.   :0.0000   Min.   :0.0000   Min.   :0.000011   Min.   :1.190
## 1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.325257   1st Qu.:1.690
## Median :0.0000   Median :0.0000   Median :0.600000   Median :2.090
## Mean   :0.1477   Mean   :0.1617   Mean   :0.565782   Mean   :1.962
```

```
## 3rd Qu.:0.0000 3rd Qu.:0.0000 3rd Qu.:0.850873 3rd Qu.:2.130
## Max. :1.0000 Max. :1.0000 Max. :0.999947 Max. :2.290
## SalePriceCH PriceDiff PctDiscMM PctDiscCH
## Min. :1.390 Min. : -0.6700 Min. :0.0000 Min. :0.00000
## 1st Qu.:1.750 1st Qu.: 0.0000 1st Qu.:0.0000 1st Qu.:0.00000
## Median :1.860 Median : 0.2300 Median :0.0000 Median :0.00000
## Mean :1.816 Mean : 0.1465 Mean :0.0593 Mean :0.02731
## 3rd Qu.:1.890 3rd Qu.: 0.3200 3rd Qu.:0.1127 3rd Qu.:0.00000
## Max. :2.090 Max. : 0.6400 Max. :0.4020 Max. :0.25269
## ListPriceDiff
## Min. :0.000
## 1st Qu.:0.140
## Median :0.240
## Mean :0.218
## 3rd Qu.:0.300
## Max. :0.440
```

```
summary(test)
```

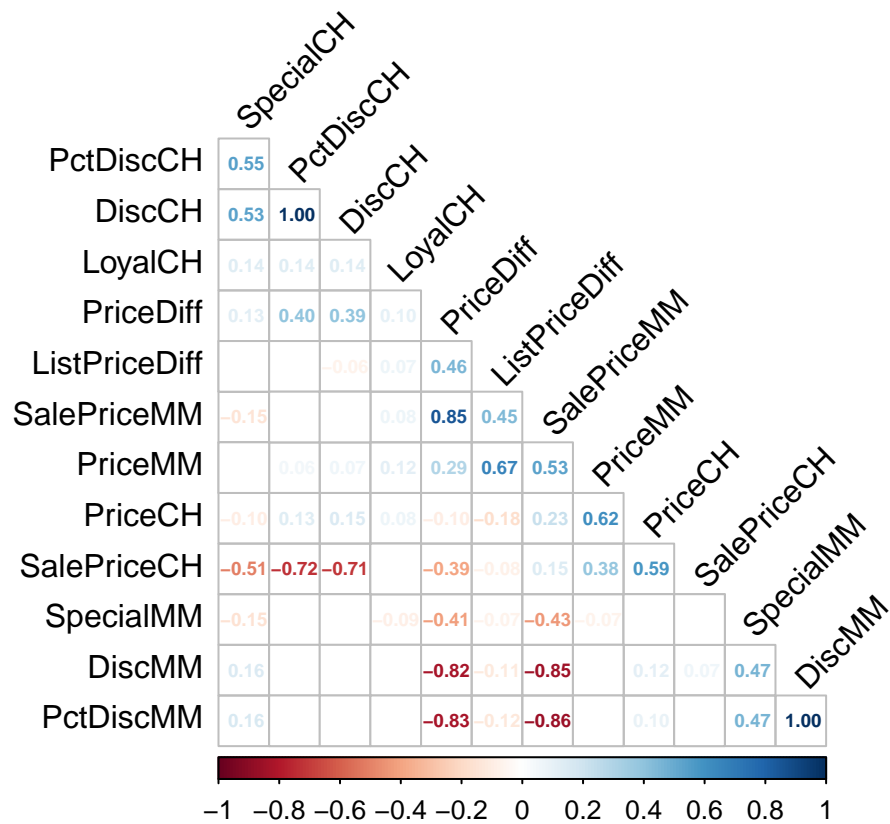
```
## Purchase PriceCH PriceMM DiscCH DiscMM
## 0:105 Min. :1.690 Min. :1.690 Min. :0.00000 Min. :0.0000
## 1:164 1st Qu.:1.790 1st Qu.:1.990 1st Qu.:0.00000 1st Qu.:0.0000
## Median :1.860 Median :2.090 Median :0.00000 Median :0.0000
## Mean :1.874 Mean :2.079 Mean :0.05167 Mean :0.1094
## 3rd Qu.:1.990 3rd Qu.:2.180 3rd Qu.:0.00000 3rd Qu.:0.2000
## Max. :2.090 Max. :2.290 Max. :0.50000 Max. :0.8000
## SpecialCH SpecialMM LoyalCH SalePriceMM
## Min. :0.0000 Min. :0.0000 Min. :0.000011 Min. :1.190
## 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.384000 1st Qu.:1.780
## Median :0.0000 Median :0.0000 Median :0.635200 Median :2.090
## Mean :0.1264 Mean :0.1413 Mean :0.595184 Mean :1.969
## 3rd Qu.:0.0000 3rd Qu.:0.0000 3rd Qu.:0.875808 3rd Qu.:2.130
## Max. :1.0000 Max. :1.0000 Max. :0.999870 Max. :2.290
## SalePriceCH PriceDiff PctDiscMM PctDiscCH
## Min. :1.390 Min. : -0.6700 Min. :0.00000 Min. :0.00000
## 1st Qu.:1.750 1st Qu.: 0.0000 1st Qu.:0.00000 1st Qu.:0.00000
## Median :1.860 Median : 0.2300 Median :0.00000 Median :0.00000
## Mean :1.823 Mean : 0.1468 Mean :0.05231 Mean :0.02709
## 3rd Qu.:1.890 3rd Qu.: 0.3000 3rd Qu.:0.09569 3rd Qu.:0.00000
## Max. :2.090 Max. : 0.6400 Max. :0.40201 Max. :0.25269
## ListPriceDiff
## Min. :0.0000
## 1st Qu.:0.1000
## Median :0.2400
## Mean :0.2045
## 3rd Qu.:0.2900
## Max. :0.4400
```

```
summary(train) #need to equalize the 0/1 split in train data set
```

```
## Purchase PriceCH PriceMM DiscCH DiscMM
## 0:312 Min. :1.690 Min. :1.690 Min. :0.00000 Min. :0.0000
## 1:489 1st Qu.:1.790 1st Qu.:1.990 1st Qu.:0.00000 1st Qu.:0.0000
## Median :1.860 Median :2.090 Median :0.00000 Median :0.0000
## Mean :1.865 Mean :2.088 Mean :0.05192 Mean :0.1281
```

```
##          3rd Qu.:1.990   3rd Qu.:2.180   3rd Qu.:0.00000   3rd Qu.:0.2400
##          Max.    :2.090   Max.    :2.290   Max.    :0.50000   Max.    :0.8000
##   SpecialCH      SpecialMM      LoyalCH      SalePriceMM
##   Min.    :0.0000   Min.    :0.0000   Min.    :0.000014   Min.    :1.19
##   1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.320000   1st Qu.:1.69
##   Median :0.0000   Median :0.0000   Median :0.585435   Median :2.09
##   Mean    :0.1548   Mean    :0.1685   Mean    :0.555908   Mean    :1.96
##   3rd Qu.:0.0000   3rd Qu.:0.0000   3rd Qu.:0.836160   3rd Qu.:2.18
##   Max.    :1.0000   Max.    :1.0000   Max.    :0.999947   Max.    :2.29
##   SalePriceCH      PriceDiff      PctDiscMM      PctDiscCH
##   Min.    :1.390   Min.    :-0.6700   Min.    :0.00000   Min.    :0.00000
##   1st Qu.:1.750   1st Qu.: 0.0000   1st Qu.:0.00000   1st Qu.:0.00000
##   Median :1.860   Median : 0.2400   Median :0.00000   Median :0.00000
##   Mean    :1.813   Mean    : 0.1464   Mean    :0.06164   Mean    :0.02739
##   3rd Qu.:1.890   3rd Qu.: 0.3200   3rd Qu.:0.11834   3rd Qu.:0.00000
##   Max.    :2.090   Max.    : 0.6400   Max.    :0.40201   Max.    :0.25269
##   ListPriceDiff
##   Min.    :0.0000
##   1st Qu.:0.1400
##   Median :0.2400
##   Mean    :0.2225
##   3rd Qu.:0.3000
##   Max.    :0.4400
```

```
corr <- cor(df[-1]) #correlogram of numeric variables, excluding outcome variable
testDf <- cor.mtest(df[-1], conf.level = 0.95) #compute significance of correlation
# Plot correlogram
corrplot(corr, p.mat = testDf$p, method = 'number', type = 'lower', insig='blank',
          addCoef.col = 'black', number.cex = 0.6, order = 'AOE', diag=FALSE, tl.srt = 45, tl.col = 'black')
```

TEST TEST TEST *Chris gearheart's test*